Basin hopping on Lennard-Jones clusters with one preset at every size

Rohit Goswami,[†,‡,¶] Ruhila Goswami,[§] Amrita Goswami,[‖] Moritz Sallermann,[‖] Sonaly Goswami,[⊥] and Debabrata Goswami[*,#,⊥]

†*SURF B.V., Amsterdam*

‡*Science Institute, University of Iceland, Reykjavik, Iceland*

¶*TurtleTech ehf., Reykjavik, Iceland*

§*Faculty of Life and Environmental Sciences, University of Iceland, Reykjavik, Iceland*

‖*Universidad Complutense de Madrid, Madrid, Spain*

⊥*Department of Chemistry, Indian Institute of Technology Kanpur, Kanpur, India*

#*Center for Bits and Atoms, Massachusetts Institute of Technology, Cambridge, Massachusetts 02139, United States*

E-mail: dgoswami@iitk.ac.in

Abstract

We describe one basin-hopping preset for Lennard-Jones clusters in which the move library, the temperature, the census, and the orbit completion stay fixed at every size. Below 98 atoms the constructor adds a diameter penalty at 0.7 of the largest pair distance, and restart on stall stays off. The novelty is that fixed core, with the depth posterior choosing the move. On the water hexamer the six rows are single draws from CREST 3.0.2, which exposes no seed, while the anneal seed fixes only start.xyz, so those rows are not a success rate: five polished gaps lie within 1×10^{-7} Eh and seed 1 lies at 1.35×10^{-7} Eh, above that line.

1 Introduction

We describe one preset that a reader can run at every cluster size. In contrast, older presets changed with the atom count. A thermodynamic model starts from a local minimum of the potential.[1] The same model treats the global minimum. As the ground state. On the Lennard-Jones pair potential the number of minima grows quickly. With the number of atoms.[2,3] A search is then judged by how many. Evaluations it spends before it reaches a minimum that is already known. Li and Scheraga introduced basin hopping for a rough. Potential-energy surface.[4] Wales and Doye made that.

Monte Carlo walk the working method for atomic clusters.[2] A move proposes a new geometry. A local minimisation returns a minimum. And the Metropolis rule at one temperature accepts or rejects. That minimum. We find the same preset at every size in the campaigns below. We report the 192-seed campaign. The preset measured here draws move arms by Thompson sampling. From a posterior on the depth of the quench. And it abandons a quench that is returning to the occupied basin.[5,6] A stall applies tabu, symmetrisation, the recorded basin trail, and the climb. The constructor installs one move library. Below 98 atoms it installs one two-phase surface, a relative cutoff of. 0.7 with weight 1, and at 98 atoms and above that list stays empty. A caller who wants lattice growth sets it on the configuration.

2 The search

2.1 One step of basin hopping

A chain stores a minimum x with energy $E(x)$. The move adds a displacement, the local minimiser returns a new minimum. y, and the chain accepts y. With the Metropolis

probability at temperature T,

$$\min\bigl(1, \exp[-(E(y) - E(x))/T]\bigr). \tag{1}$$

T is 0.8 in the reduced units where the pair well. Depth is 1. Wales and Doye use that reduced temperature for Lennard-Jones. Basin hopping.[2] The Lennard-Jones preset stores the value as `TEMPERATURE`. One temperature is used at every size. The hit that scores a run is the first converged evaluation whose. Energy lies below the published minimum by less than 0.001. A miss is a run that reaches its call ceiling without such. An evaluation. The penalised average runtime (PAR-2) is the call count of a hit. And twice the cap for a miss. The hit is read from the trace of converged improvements. The trace stores no stop threshold.

Algorithm 1 lists one chain of the recommended preset.

Algorithm 1 One chain of the recommended preset.

1: Draw a non-overlapping start from the seed and quench it.
2: **while** the call budget remains **do**
3: Draw a move from the library. And quench the displaced geometry to a minimum y.
4: Replace x by y. With the probability in equation 1.
5: When y is a new basin. Quench the symmetry images of its core.
6: **end while**

2.2 What the preset turns on

The recommended preset starts from the Wales and Doye displacement. And then selects one move library. Thompson sampling draws each move from a normal-gamma posterior. On the depth reached by earlier quenches of that move.[5] Below 98 atoms the constructor sets the surface. List to one diameter penalty. By calling `TwoPhase::relative` with. 0.7 and 1. The cutoff equals 0.7 times the largest pair distance. That distance is read on the structure that enters the relaxation. The penalty weight equals 1. The surface portfolio then draws that penalty or the plain surface, using the same depth posterior. The surface

list stays empty at 98 atoms and above. The relaxation then uses the plain surface. The lattice-regrowth library builds a truncated octahedron. A decahedron, or an icosahedron from a named template. The global minimum changes with size.[2,7] Face-centred cubic. Is the minimum at 38 atoms, icosahedral at 55, and Marks decahedral at. 75. Installing that library only at 38 atoms would search for a structure. The count has already named. The preset therefore keeps the lean burst library at every size, including 38. Doye showed that the Locatelli-Schoen compression. A term that pulls every atom toward the centroid. Is a change of the potential.[8,9] The installed surface is a diameter. Penalty, and its centroid weight stays at 0. Restart on stall stays at its default, off. So a stalled walker is not redrawn from random coordinates. A stall still applies tabu, symmetrisation, the recorded basin trail, and the climb. Orbit completion quenches the images of a symmetric core when a new. Basin appears. A continuous-symmetry projection every ten quenches. And a reseed after five hops on an icosahedral plateau. Follow the defaults Oakley, Johnston and Wales published for the Wales-group basin-hopping program. (GMIN).[10] Those two numbers are applied at every size. Figure 1 stacks the potential, the hop. And the bank that shares the minima. The figure is drawn from the search, and no campaign count enters it.

3 Statistical framework

A cooling schedule sets the step. And the temperature of one chain. A population of chains also decides which replicas carry on. The census of packings then decides when the next minimization belongs. To a different packing. We resample replicas on within-epoch ranks, and we close the census. With a Good-Turing estimate of the unseen mass.[12]

3.1 Annealed-flow kernel

As the search cools, the replicas follow a sequence of distributions. that weight lower energies more heavily. Del Moral, Doucet, and Jasra write that sequence as the equation below.[13]

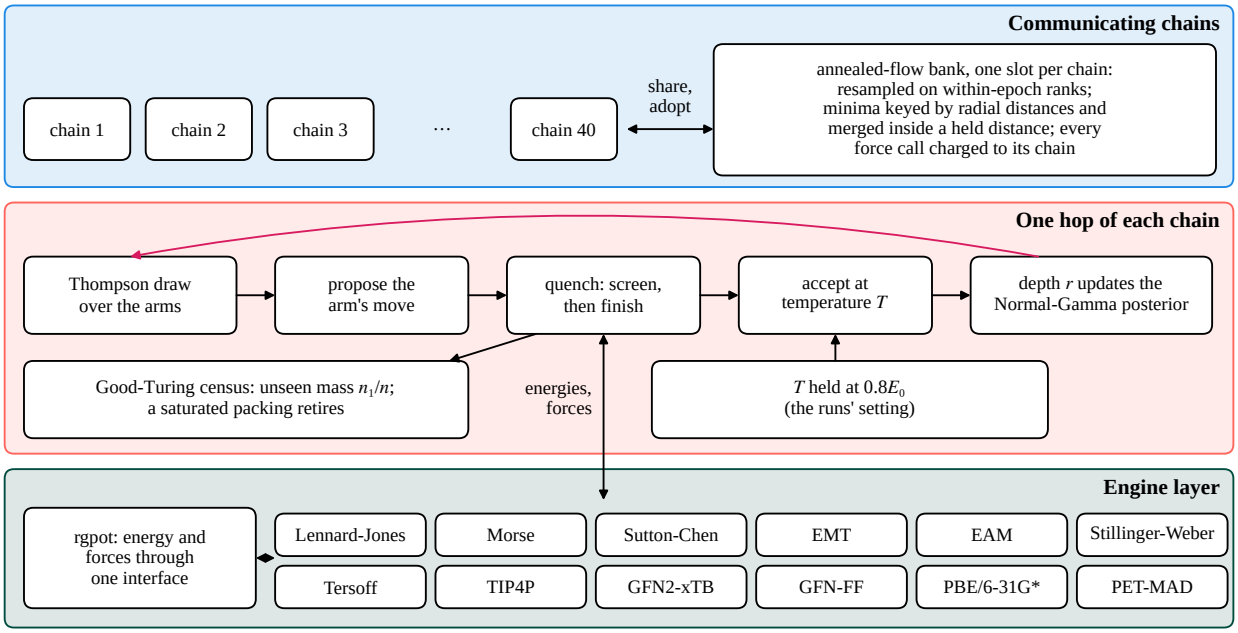


Figure 1: The stack of the search. The chains share minima through the annealed-flow bank. In one hop the Thompson draw picks the arm, the quench is screened. And the acceptance test runs at the held temperature $T = 0.8$. The depth of the quench updates the posterior of the arm that played, and a saturated packing retires. Every potential reaches the hop through rgpot.[11]

$$\eta_t(\mathrm{d}x) \propto \exp(-\beta_t E(x))\, \lambda(\mathrm{d}x), \tag{2}$$

Beta increases along the sequence. They move from one target to the next by reweighting. They then apply a move kernel.

$$\eta_{t+1} = \Psi_{G_t}(\eta_t)\, M_{t+1}. \tag{3}$$

Cluster energies share no scale from one size to the next. The weight is therefore the within-epoch rank. Tied replicas share their average position. Positions run from 0 to $R-1$. A tied block runs from b inclusive to e exclusive. The block has the rank below when $R \geq 2$.

$$r = \frac{b + e - 1}{2(R-1)}. \tag{4}$$

A lower energy receives a lower rank. The energy rank r_E, the novelty rank r_N, the scarcity rank. r_S, and the uncertainty rank r_U combine as

$$\log w_i = -r_{E,i} + 0.8\, r_{N,i} + 0.6\, r_{S,i} + 0.4\, r_{U,i}. \tag{5}$$

Offspring go to replicas that are low in energy, rare as packings. And uncertain as basins.

3.2 Occupancy and the temperature of the hop

Those packings and basins set the occupancy of the hop. The Good-Turing estimate of the unseen mass closes that census.[12] A saturated packing. Then retires. The next minimization is spent on a packing the census has barely seen. The temperature of the recommended hop stays at the Wales and Doye value 0.8 in reduced units.[2] A caller who wants a schedule that tracks the entropy. Of the occupied states sets it on the configuration. The preset itself keeps one temperature at every size.

4 Communicating chains

Many searches already run many chains. We call the communication correct when five conditions hold. Each chain draws its next move from a posterior on the depth. That move reaches. A packing retires on the Good-Turing estimate of its unseen mass. Minima are shared under one radial merge distance. Every force evaluation is charged to the chain that spent it. So chains are compared at equal budgets. The receiver can act on a shared minimum by quenching it. And entering it in the acceptance test.

Parallel tempering (PT) exchanges states along a ladder of temperatures.[14,15] Population annealing resamples a population by its Boltzmann weight.[16,17] Conformational space annealing (CSA) keeps a bank of minima.[18] CSA also anneals. The distance cutoff. Deaven and Ho mate stored geometries by cut and splice.[19] The Birmingham pool algorithms do the same.[20] Basin hopping (BH) and minima hopping (MH) run a single walker.[2,6] The conformer-rotamer ensemble sampling. Tool (CREST) samples a molecular ensemble by metadynamics.[21] The bank in this section adds. The depth posterior and the unseen-mass census to that list.

5 Information theory in the hop

A move that has stopped finding depth can still take. Over the next trial. The recommended hop guards against that takeover. With a posterior on the depth each move reaches.

5.1 Depth posterior

On a deep cluster few of the available moves improve the incumbent. We draw the next arm by Thompson sampling from a normal-gamma posterior. On the depth earlier quenches of that arm reached.[5] The depth that updates the posterior is the drop. From the energy before the quench to the energy after it. The same acceptance test meets every arm. With the well-tempered bias at zero. That test is equation 1.[22] Figure 2 draws. The arms and the

capped burst length. The figure is drawn from the search, and no campaign count enters it.

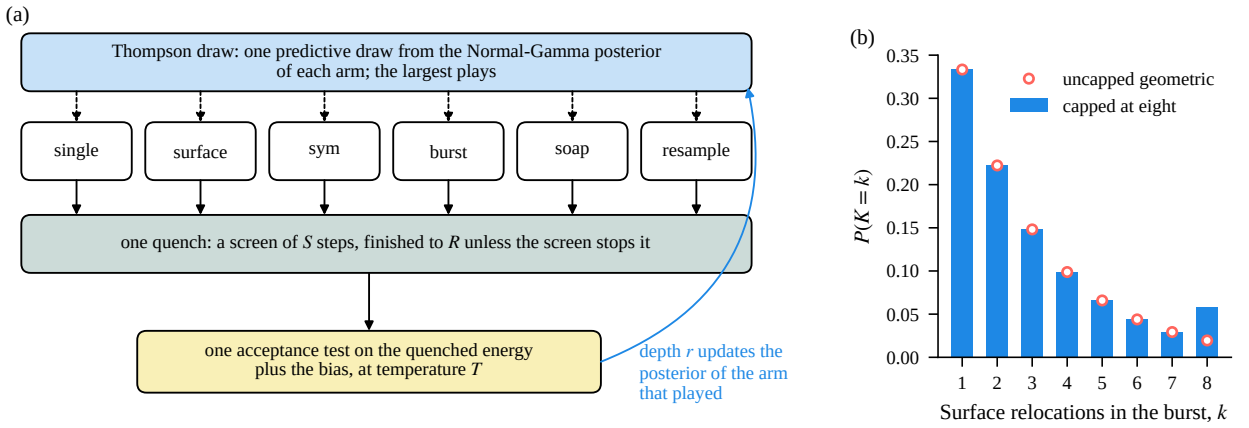


Figure 2: The figure shows the arms of the recommended hop. And the law of its burst. The Thompson draw picks one arm. Whichever arm plays, its proposal is quenched once. That proposal meets one acceptance test. The depth then updates the posterior. The burst length is capped at eight relocations.

5.2 Symmetrisation

The burst drawn above is one arm of the hop. After that quench. The search projects the structure toward a more symmetric neighbour. A global minimum is often more symmetric than the minima around it. On LJ_{75} the global minimum is a Marks decahedron.[7,23] Most. Low minima around that decahedron are icosahedral. Oakley, Johnston, and Wales symmetrise the structure in GMIN.[10] One route completes the orbits of a symmetric core. Another route uses a continuous symmetry measure. The preset applies that projection every ten quenches at every size.

6 Results

6.1 The structures

Figure 3 draws the archived structures of the benchmark. The Lennard-Jones geometries are the reference minima of the Cambridge. Cluster Database, and the TIP4P geometries

are the minima of Wales and. Hodges.[24–26] The hexamers are minima of the hop at the GFN2-xTB tight-binding model.[27] The Supporting Information. Draws the full structure plates.

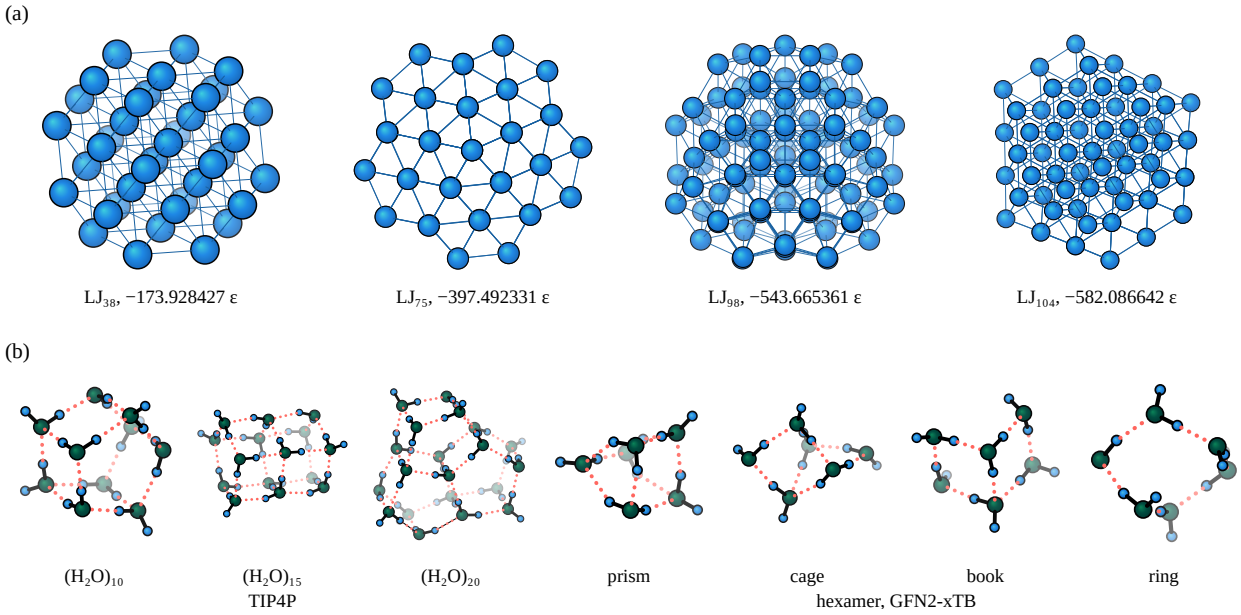


Figure 3: Structures of the benchmark. (a) The global minima of LJ_{38}, LJ_{75}, LJ_{98} and LJ_{104}. (b) The TIP4P minima of 10, 15 and 20 molecules. And the lowest GFN2-xTB hexamer of each class.

The LJ_{75} seeds 0 through 199 run serially at a cap of. 4×10^6 force calls. No hit count for those seeds is in the sealed record. Figure 4 draws the encounter curves of the earlier. Sealed LJ_{75} campaign, which covered 200 seeds. The old-method baseline on LJ_{104} is measure v4.14. Its six receipts, seeds 2000 through 2005, record 0 of 6 hits, and the best of those receipts is −581.80180238 on seed 2005, against the reference −582.086642.

6.2 Rates and the water hexamer

Figure 5 maps the minima one seed of basin hopping and of the recommended hop accepted at LJ_{38}, LJ_{75} and. LJ_{98}. At LJ_{38} the recommended hop reached the truncated octahedron. The LJ_{75} minima on that map reach the icosahedral competitor at. −396.282249, which is 1.210 above the decahedral global minimum. At −397.492331. The LJ_{98} minima stop 0.022

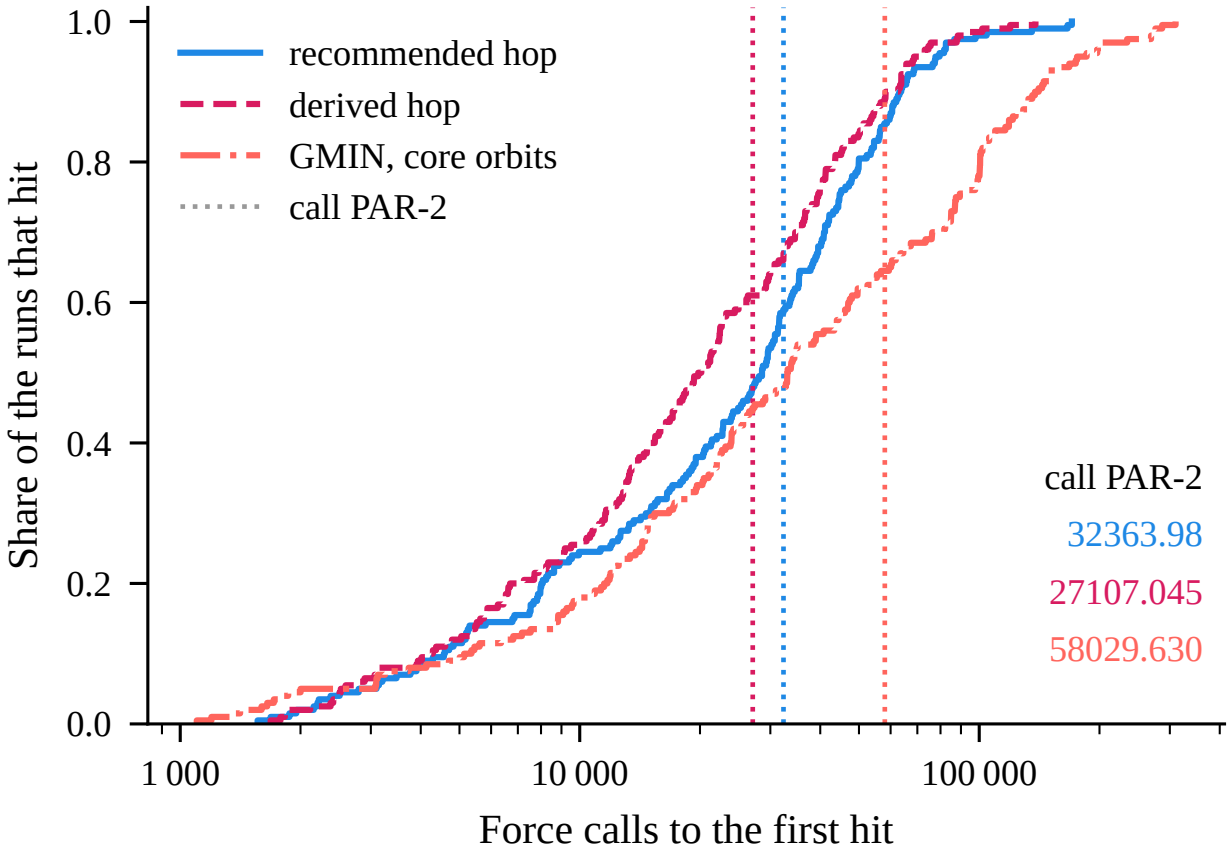


Figure 4: Encounter curves of main's earlier sealed LJ_{75} campaign. The panel prints that run's PAR-2 values. The run used an empty surface list and restart on stall. The arms are the recommended hop, the derived hop. And the Wales-group basin-hopping program (GMIN). The plot covers 200 seeds.

above the tetrahedral global. Minimum at -543.665361. Figure 6 joins the LJ_{38} minima where. They first connect among their nearest neighbours. In the coordination descriptor. A search records no transition states. So no join of the tree is a barrier. Figure 7 shows a landfold panel for each sealed Lennard-Jones size and for the TIP4P hexamer.

6.3 Earlier 192-seed hexamer campaign

On the water hexamer the sealed campaign runs four arms. Each on seeds 0 to 191 at a budget of 4000 force calls.[21,27] The default arm is the hop. The budget-fair polish is an opt-in mode that charges an external. Polish against the same budget. The third arm is CREST. And the fourth runs that polish after CREST inside the same budget. A run hits when the iterative rotations and assignments alignment (IRA). matches a structure to a Malloum prism, cage, book. Or cyclic chair reoptimised at GFN2-xTB.[28–30] The default hop hit 121 of 192 seeds. Its call PAR-2 median is 4000 and its type-7 interquartile range is. 4000. CREST hit 97 of 192 seeds. Its call PAR-2 median is 4000 and its type-7 interquartile range is. 4000. Paired by seed, the exact McNemar test of the discordant pairs gives. $p = 0.0264$.[31] The budget-fair polish hit 148 of 192. Its call PAR-2 median is 4000 and its type-7 interquartile range is 0. CREST plus that polish hit 114 of 192. Its call PAR-2 median is 3560 and its type-7 interquartile range is. 4546.25. The exact McNemar p of the polish against CREST is. 1.37×10^{-6}. The exact McNemar p against CREST plus the polish is 0.000374. Figure 8 draws the four arms.

6.4 Water hexamer against CREST

CREST here is conda-forge crest 3.0.2. Its metadynamics is stochastic for a fixed start geometry. This build exposes no seed. The anneal seed fixes only start.xyz. CREST's cost is the energy and gradient call count. The matched comparison. And the 30-start statistics are not in this article. A row succeeds when the run exits cleanly. The polished energy is finite. And the maximum force is at or below 0.001 Eh/bohr. Every final is polished with

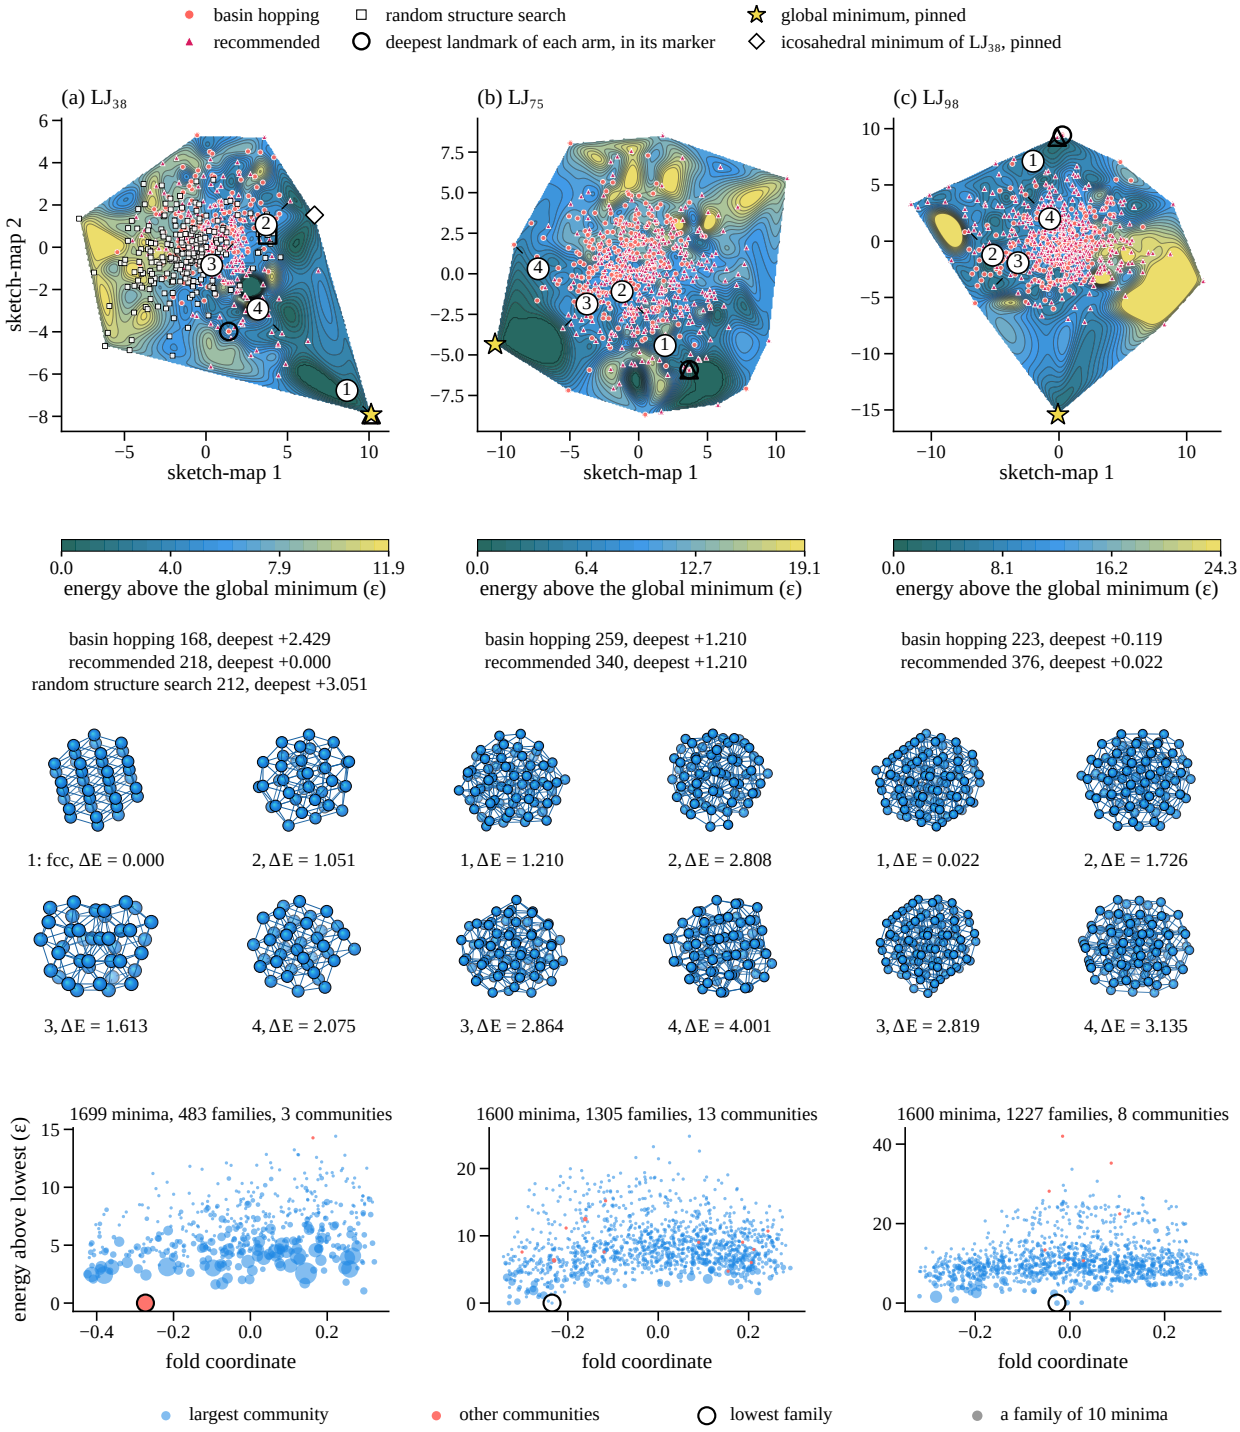


Figure 5: Maps of the minima one seed accepted at LJ_{38}, LJ_{75}, and LJ_{98}. Top, the sketch-map. Middle, four deep landmarks. Bottom, the landfold map of the same minima. Sky in the largest community and coral in the others.

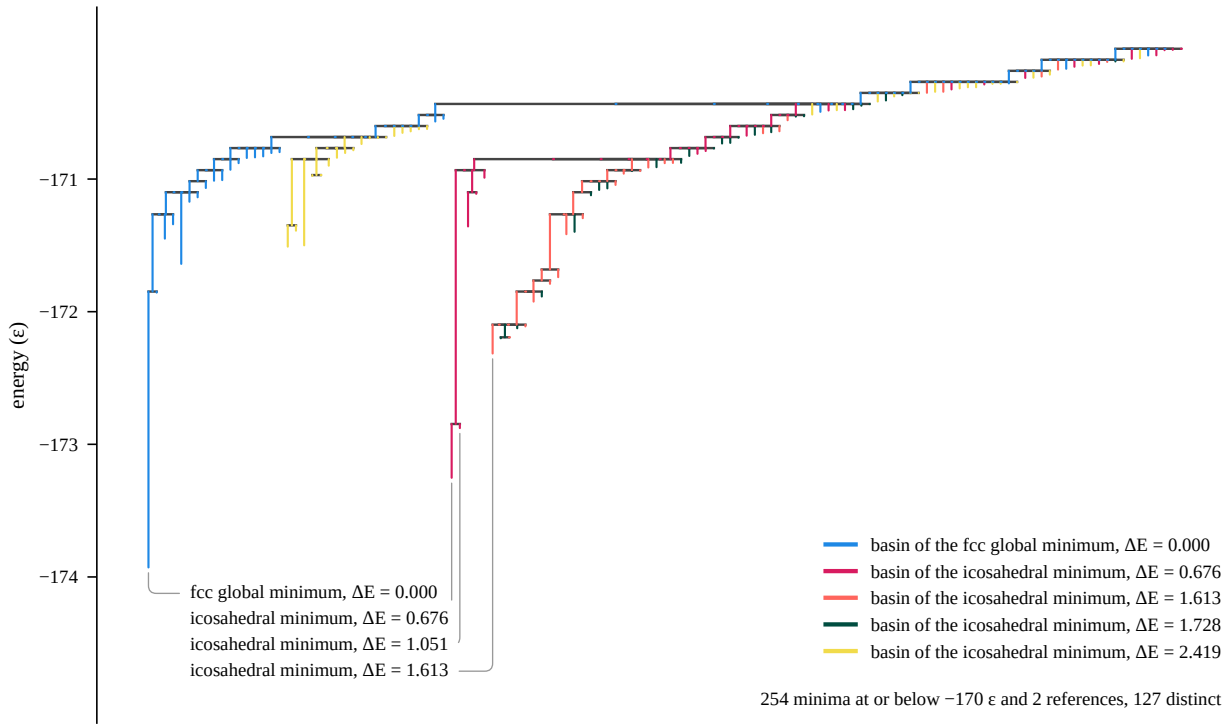


Figure 6: Approximate disconnectivity graph of LJ_{38}. Minima closer than 0.15 in the coordination descriptor are merged. The stems are coloured by five spectral basins. No join is a barrier.

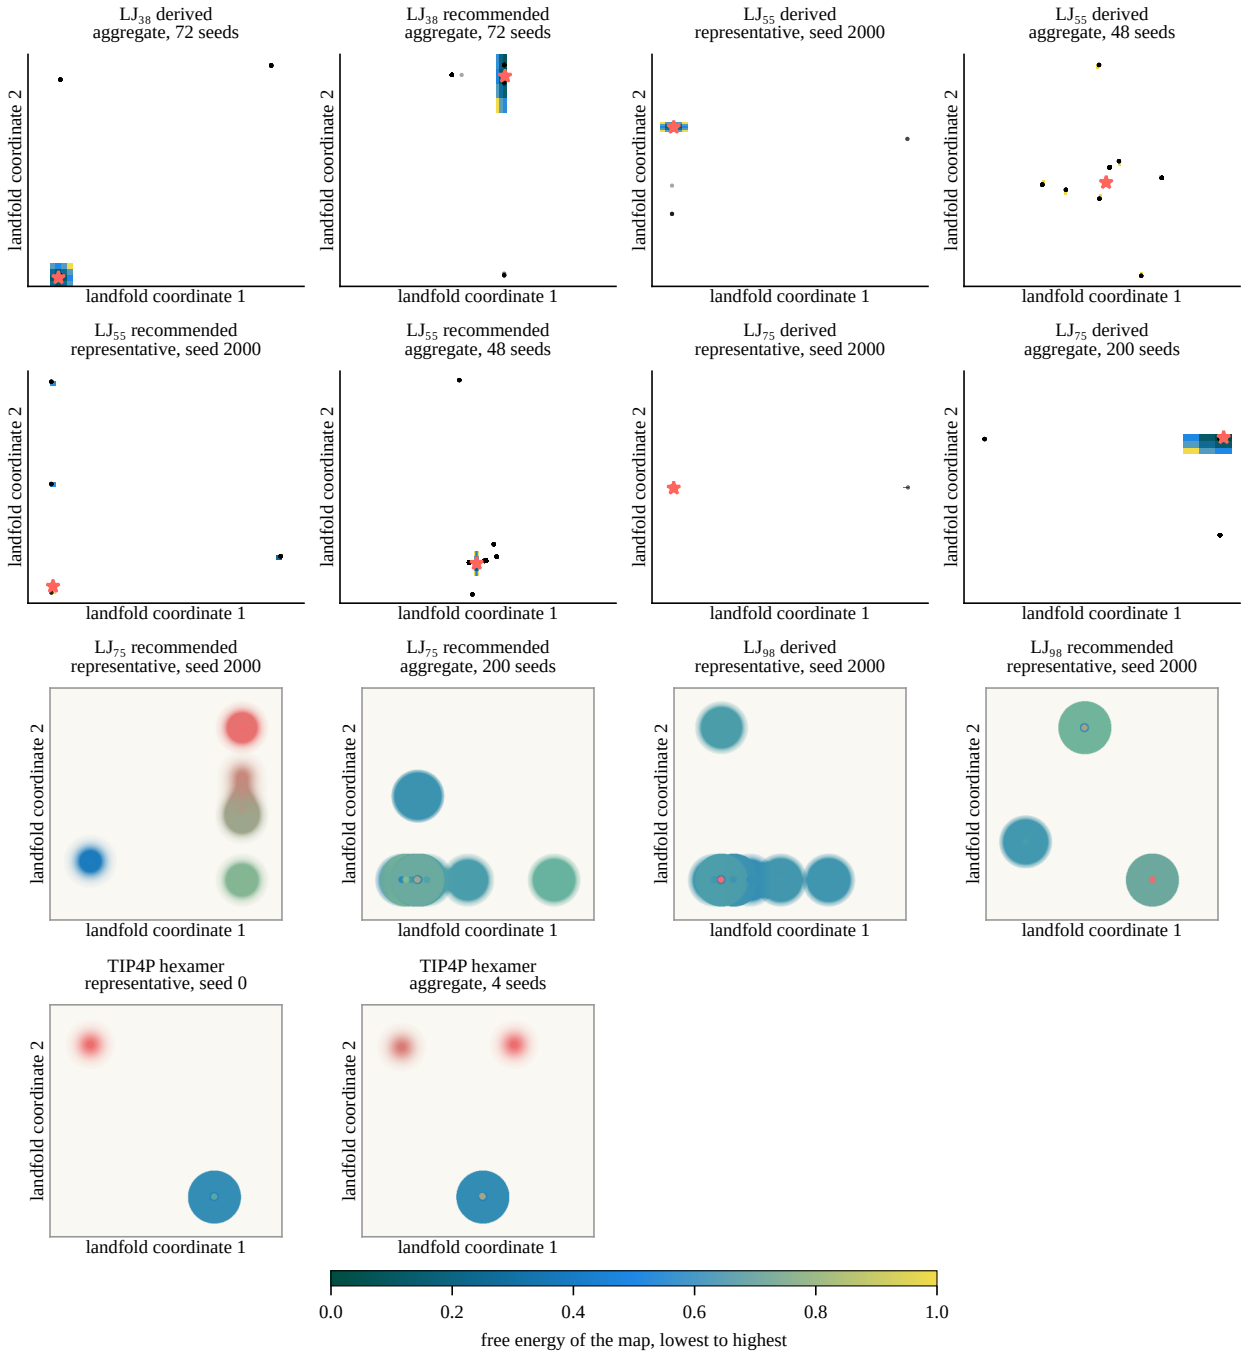


Figure 7: Landfold maps of the sealed Lennard-Jones and TIP4P hexamer runs. The Lennard-Jones panels use main's preset, including an LJ_{38}. derived panel and an LJ_{75} aggregate of 200 seeds. Each panel is one sealed map. Coloured from the lowest free energy to the highest. With the minimum marked.

xtb at GFN2, the vtight setting, accuracy. 0.01, and an iteration limit of 500. The vtight energy threshold is 1×10^{-7} Eh.

Table 1 lists two CREST counters for each seed. The old counter sums the dynamics steps and the optimization steps. CREST prints a second total, "Total number of energy+grad calls". The old counter undercounted CREST on every seed.

Table 1: CREST effort on the water hexamer. The old counter is dynamics steps plus optimization steps. The energy and gradient column is CREST's own summary line.

Seed	Old counter	Energy+gradient
0	32112	151074
1	50026	191132
2	32155	142110
3	46062	172468
4	50063	157643
5	32170	209931

The 4000-call arm spends 4000 evaluations in one setting. The driver-count budget is the anneal arm whose cap equals each seed's. Old counter. The driver-count budget runs from 32112 to 50063 evaluations. CREST's energy and gradient totals run from 142110 to 209931. The comparison is the 4000-call arm against those energy. And gradient calls. The driver-count budget is not equal effort. Because the old counter undercounted CREST.

Table 2 gives the polished gap on each seed. The gap is the anneal energy minus the CREST energy. A negative value means the anneal energy sits lower. Five seeds lie within 1×10^{-7} Eh on both gap ledgers. Seed 1 exceeds that line, at 1.3509×10^{-7} Eh on the. 4000-call ledger. And 1.3505×10^{-7} Eh on the driver-count ledger. The label-free root-mean-square deviation of each polished pair. Is at most 0.00212 Å.

The logs name the processor Intel Xeon Processor, family 6, model 207. Each CREST run selected one thread. The hexamer wall-time medians, from `logs/hexamer-walls.txt`, are 544.0 s for CREST, 18.8 s for the 4000-call arm. And 195.5 s for the driver-count budget.

The CREST polish converged in one iteration on all 18 finals. No CREST structure dissociated.

Table 2: Polished gap and label-free deviation on the water hexamer. The gap is anneal minus CREST, in hartree. The deviation is the label-free root-mean-square deviation. In angstroms.

Seed	Gap, 4000-call ledger	Gap, driver-count ledger	Deviation, 4000-call ledger	Deviation, driver-count ledger
0	-1.1214×10^{-8}	-1.1210×10^{-8}	0.001004	0.000987
1	-1.3509×10^{-7}	-1.3505×10^{-7}	0.002023	0.001964
2	-2.5460×10^{-9}	-2.5850×10^{-9}	0.000532	0.000477
3	-4.1510×10^{-9}	-4.1430×10^{-9}	0.000322	0.000308
4	-6.8980×10^{-8}	-6.9022×10^{-8}	0.002116	0.002068
5	-2.1900×10^{-9}	-2.2000×10^{-9}	0.000157	0.000134

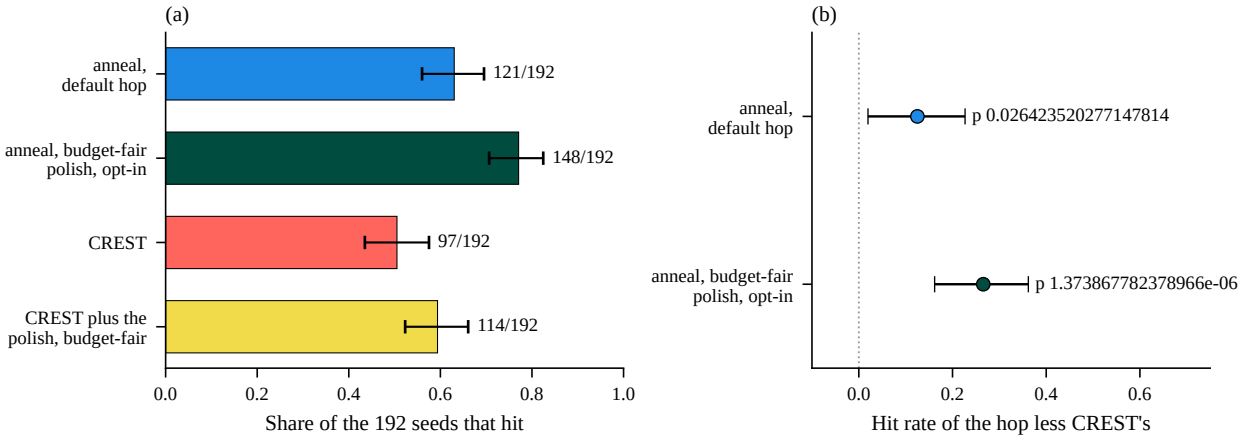


Figure 8: The water hexamer against CREST, sealed campaign. (a) The share of 192 seeds that each arm hit, with the Wilson interval.[32] (b) For each hop arm paired with CREST, the difference in hit rate, with the Newcombe interval. And the exact McNemar p.[31,33]

The water 16-mer count is separate open work. Apart from the hexamer campaign. This comparison leaves out the Atomic Simulation Environment. The atomistic global optimization framework (AGOX). And the Wales-group basin-hopping program (GMIN).

7 Conclusions

On the water hexamer the sealed 192-seed hop hit 121 of 192 seeds. In contrast, CREST hit 97 of 192 seeds. The budget-fair polish hit 148 of 192, and CREST followed. By that polish hit 114 of 192. Five of the six hexamer seeds lie within 1×10^{-7} Eh on both. Gap ledgers. Seed 1 lies above that line. No hit count for the LJ_{75} seeds 0 through 199 is in the sealed record.

Author contributions

Author contributions follow the CRediT taxonomy. Rohit Goswami contributed conceptualization, methodology, software, formal analysis, and the original draft. Ruhila Goswami contributed software and visualization. Amrita Goswami contributed methodology and edited the draft. Moritz Sallermann contributed software and methodology. And he edited the draft. Sonaly Goswami contributed validation and checked the draft. Debabrata Goswami contributed supervision and oversaw the draft.

Supporting Information

Board constants, reference-structure plates, sketch-map panels, and landfold panels of the Lennard-Jones maps (PDF)

Acknowledgments

Rohit Goswami thanks his pets Arí, Yoda, Jude, Crystee, Thor, and. Loki, and the garden cats. In memoriam TuiTui, Snoopy, and Cosmo. He thanks his cactus Ami and his in-laws Nagaselvamani and. Sankaranarayan. Ruhila Goswami set the figure colours. Generative AI tools assisted with drafting and editing the text. The authors reviewed and edited all content. And take full responsibility for it.

Funding

This research received no specific grant. No funding agency in the public, commercial. Or not-for-profit sectors supported it.

Notes

The authors declare no competing financial interest.

Data Availability Statement

The data supporting the conclusions are in this article, the Supporting Information, and the public repositories named here. The source is the repository `https://github.com/HaoZeke/anneal`. under the Massachusetts Institute of Technology (MIT) licence. Documentation is at `https://anneal.rgoswami.me`. The archive of the software line is the Zenodo concept record. `https://doi.org/10.5281/zenodo.10672746`. The Zenodo digital object identifier (DOI) for this submission is not assigned. The newest released tag is v0.11.1. The crate records version 0.12.0. The release tag for version 0.12.0 is not assigned. The copper numbers come from the slab runs. Commit aef5cf0b2068c4988509066a881eb68682ad9290 produced them. The rerun best energies run from -415.971521 eV to -415.971161 eV. The

hit tolerance is 0.001 eV. The script is `scripts/elja_hq_sci_molslab_bank.sh`. The lines are 118 and 119. TARGET_TOL there is 1e-3. The per-seed gaps run from 8e-6 eV to 3.68e-4 eV above -415.971529 eV. No seed reached the reference within 1e-6 eV. These are not exact hits. Eight seeds lie inside the 0.001 eV tolerance. The `set_relative_0.01` line records the 0.01 eV rate. That line is a separate figure. The library pins rgpot 25c759da82eececa8897e4634046fab27409c81a. The library sha256 follows.

d399ef1f675e97c1f5e401a987d6a9df1bc30b3a37ba6586044b66d44be8399d.

References

(1) Wales, D. J. *Energy Landscapes*; Cambridge University Press: Cambridge, U.K., 2003.

(2) Wales, D. J.; Doye, J. P. K. Global optimization by basin-hopping and the lowest energy structures of Lennard-Jones clusters containing up to 110 atoms. *J. Phys. Chem. A* **1997**, *101*, 5111–5116.

(3) Jones, J. E. On the determination of molecular fields. II. From the equation of state of a gas. *Proc. R. Soc. London, Ser. A* **1924**, *106*, 463–477.

(4) Li, Z.; Scheraga, H. A. Monte Carlo-minimization approach to the multiple-minima problem in protein folding. *Proc. Natl. Acad. Sci. U.S.A.* **1987**, *84*, 6611–6615.

(5) Thompson, W. R. On the likelihood that one unknown probability exceeds another in view of the evidence of two samples. *Biometrika* **1933**, *25*, 285–294.

(6) Goedecker, S. Minima hopping: An efficient search method for the global minimum of the potential energy surface of complex molecular systems. *J. Chem. Phys.* **2004**, *120*, 9911–9917.

(7) Marks, L. D. Surface Structure and Energetics of Multiply Twinned Particles. *Philosophical Magazine A* **1984**, *49*, 81–93.

(8) Doye, J. P. K. The effect of compression on the global optimization of atomic clusters. *Phys. Rev. E* **2000**, *62*, 8753–8761.

(9) Locatelli, M.; Schoen, F. Fast global optimization of difficult Lennard-Jones clusters. *Comput. Optim. Appl.* **2002**, *21*, 55–70.

(10) Oakley, M. T.; Johnston, R. L.; Wales, D. J. Symmetrisation schemes for global optimisation of atomic clusters. *Phys. Chem. Chem. Phys.* **2013**, *15*, 3965–3976.

(11) Goswami, R. OmniPotent RPC rgpot. Zenodo, version 3.5.1, 2026; `https://github.com/OmniPotentRPC/rgpot`.

(12) Good, I. J. The Population Frequencies of Species and the Estimation of Population Parameters. *Biometrika* **1953**, *40*, 237–264.

(13) Del Moral, P.; Doucet, A.; Jasra, A. Sequential Monte Carlo Samplers. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* **2006**, *68*, 411–436.

(14) Swendsen, R. H.; Wang, J.-S. Replica Monte Carlo Simulation of Spin-Glasses. *Physical Review Letters* **1986**, *57*, 2607–2609.

(15) Hukushima, K.; Nemoto, K. Exchange Monte Carlo Method and Application to Spin Glass Simulations. *Journal of the Physical Society of Japan* **1996**, *65*, 1604–1608.

(16) Hukushima, K.; Iba, Y. Population Annealing and Its Application to a Spin Glass. The Monte Carlo Method in the Physical Sciences: Celebrating the 50th Anniversary of the Metropolis Algorithm. 2003; pp 200–206.

(17) Machta, J. Population annealing with weighted averages: A Monte Carlo method for rough free-energy landscapes. *Physical Review E* **2010**, *82*, 026704.

(18) Lee, J.; Lee, I.-H.; Lee, J. Unbiased Global Optimization of Lennard-Jones Clusters for $N \leq 201$ Using the Conformational Space Annealing Method. *Physical Review Letters* **2003**, *91*, 080201.

(19) Deaven, D. M.; Ho, K. M. Molecular geometry optimization with a genetic algorithm. *Phys. Rev. Lett.* **1995**, *75*, 288–291.

(20) Shayeghi, A.; Götz, D.; Davis, J. B. A.; Schäfer, R.; Johnston, R. L. Pool-BCGA: a parallelised generation-free genetic algorithm for the ab initio global optimisation of nanoalloy clusters. *Physical Chemistry Chemical Physics* **2015**, *17*, 2104–2112.

(21) Pracht, P.; Grimme, S.; Bannwarth, C.; Bohle, F.; Ehlert, S.; Feldmann, G.; Gorges, J.; Müller, M.; Neudecker, T.; Plett, C.; Spicher, S.; Steinbach, P.; Wesołowski, P. A.; Zeller, F. CREST, A program for the exploration of low-energy molecular chemical space. *The Journal of Chemical Physics* **2024**, *160*, 114110.

(22) Barducci, A.; Bussi, G.; Parrinello, M. Well-Tempered Metadynamics: A Smoothly Converging and Tunable Free-Energy Method. *Physical Review Letters* **2008**, *100*, 020603.

(23) Doye, J. P. K.; Miller, M. A.; Wales, D. J. The double-funnel energy landscape of the 38-atom Lennard-Jones cluster. *The Journal of Chemical Physics* **1999**, *110*, 6896–6906.

(24) Wales, D. J.; Doye, J. P. K.; Dullweber, A.; Hodges, M. P.; Naumkin, F. Y.; Calvo, F.; Hernández-Rojas, J.; Middleton, T. F. The Cambridge Cluster Database. `https://www-wales.ch.cam.ac.uk/CCD.html`, Accessed 25 October 2026.

(25) Wales, D. J.; Hodges, M. P. Global minima of water clusters $(H_2O)_n$, $n \leq 21$, described by an empirical potential. *Chemical Physics Letters* **1998**, *286*, 65–72.

(26) Jorgensen, W. L.; Chandrasekhar, J.; Madura, J. D.; Impey, R. W.; Klein, M. L. Comparison of simple potential functions for simulating liquid water. *J. Chem. Phys.* **1983**, *79*, 926–935.

(27) Bannwarth, C.; Ehlert, S.; Grimme, S. GFN2-xTB, An Accurate and Broadly Parametrized Self-Consistent Tight-Binding Quantum Chemical Method with Multipole Electrostatics and Density-Dependent Dispersion Contributions. *Journal of Chemical Theory and Computation* **2019**, *15*, 1652–1671.

(28) Gunde, M.; Salles, N.; Hémeryck, A.; Martin-Samos, L. IRA: A Shape Matching Approach for Recognition and Comparison of Generic Atomic Patterns. *Journal of Chemical Information and Modeling* **2021**, *61*, 5446–5457.

(29) Malloum, A.; Fifen, J. J.; Dhaouadi, Z.; Nana Engo, S. G.; Conradie, J. Structures, Relative Stability and Binding Energies of Neutral Water Clusters, $(H_2O)_{2-30}$. *New Journal of Chemistry* **2019**, *43*, 13020–13037.

(30) Pérez, C.; Muckle, M. T.; Zaleski, D. P.; Seifert, N. A.; Temelso, B.; Shields, G. C.; Kisiel, Z.; Pate, B. H. Structures of Cage, Prism, and Book Isomers of Water Hexamer from Broadband Rotational Spectroscopy. *Science* **2012**, *336*, 897–901.

(31) McNemar, Q. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika* **1947**, *12*, 153–157.

(32) Wilson, E. B. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association* **1927**, *22*, 209–212.

(33) Newcombe, R. G. Improved confidence intervals for the difference between binomial proportions based on paired data. *Statistics in Medicine* **1998**, *17*, 2635–2650.

For Table of Contents Only

